

The code (for reference)

```
BLOCKED_TERMS = {'suicide', 'bomb', 'explode', 'terror', 'kill', 'illegal'}
```

```
def contains_blocked(text: str) -> bool:
```

```
    if not text: return False
```

```
    t = text.lower()
```

```
    return any(term in t for term in BLOCKED_TERMS)
```

```
def sanitize_output(text: str) -> str:
```

```
    if not text: return text
```

```
    out = text
```

```
    for term in BLOCKED_TERMS:
```

```
        out = out.replace(term, '[redacted]')
```

```
    return out
```

Step-by-step execution flow

1) When the module loads

1. Python evaluates the right-hand side and creates a **set** named `BLOCKED_TERMS` containing the six strings:
`{'suicide', 'bomb', 'explode', 'terror', 'kill', 'illegal'}`.
 - A set gives $O(1)$ average membership tests and ensures each term is unique.
-

2) Calling `contains_blocked(text)`

Suppose you call `contains_blocked(input_text)`.

2.1. Check for empty input

- `if not text: return False` — if `text` is `None` or an empty string (`""`) the function returns `False` immediately. No further work is done.

2.2. Lowercase conversion

- `t = text.lower()` — convert the entire input to lowercase so matching is case-insensitive.
Example: `"Bomb"` → `"bomb"`.

2.3. Search blocked terms

- `return any(term in t for term in BLOCKED_TERMS)`

- This creates a generator that, for each term in `BLOCKED_TERMS`, checks whether that substring appears anywhere inside `t`.
- `any(...)` short-circuits on the first `True`. If any blocked term is found, the function returns `True`; otherwise `False`.

Examples

- `contains_blocked("I will kill the spider")` → `True` (because "kill" is present)
- `contains_blocked("This is illegal activity")` → `True`
- `contains_blocked("I like sunshine")` → `False`
- `contains_blocked("")` or `None` → `False` (handled by the early return)

Edge note: This is a *substring* check — it matches inside other words (e.g., "skill" contains "kill"), which may or may not be desirable.

3) Calling `sanitize_output(text)`

Suppose you call `sanitize_output(input_text)`.

3.1. Check for empty input

- if not `text`: return `text` — if `text` is falsy (`None` or `""`), it returns it untouched.

3.2. Copy text to out

- `out = text` — working copy so the original `text` variable is not modified (strings are immutable anyway).

3.3. Loop and replace

- for `term` in `BLOCKED_TERMS`: iterates over the set (order arbitrary).
- `out = out.replace(term, '[redacted]')` replaces **all occurrences** of the exact substring `term` with `"[redacted]"`. This replacement is case-sensitive — it replaces only exact-match substrings, so if the original text had "Bomb" (capital B), `.replace("bomb", ...)` will **not** replace it unless you pre-normalize or handle case.

3.4. Return sanitized string

- After the loop finishes for all terms, return `out`.

Examples

- `sanitize_output("He built a bomb")` → `"He built a [redacted]"`
- `sanitize_output("I might KILL him")` → `"I might KILL him"` — **note:** uppercase "KILL" remains because `.replace("kill", ...)` won't match "KILL".
- `sanitize_output("")` → `""`

Important behavioral notes

- Because `BLOCKED_TERMS` is iterated in arbitrary set order, the sequence of replacements is non-deterministic across runs (not usually a huge deal for distinct terms but can matter if one term is substring of another).
 - Replacement is blind — no word-boundary detection, so "skill" becomes "s[redacted]" if "kill" is a blocked term (because "kill" occurs inside "skill").
 - Case-sensitivity mismatch between `contains_blocked` (lowercasing, case-insensitive) and `sanitize_output` (case-sensitive) can lead to surprising results.
-

Example run (full)

Input: "He tried to explode the old kiln but failed — illegal act."

`contains_blocked(...)`:

- Lowercases text and finds "explode" and "illegal" → returns True.

`sanitize_output(...)`:

- Replaces occurrences of "explode" and "illegal" (only exact lowercase matches). If input used exact lowercase, both replace — otherwise some may remain.
-

Common pitfalls & edge cases

- **Partial-word matches:** "skill" contains "kill" → could incorrectly redact part of a benign word.
- **Case sensitivity mismatch:** `contains_blocked` is case-insensitive; `sanitize_output` is not. That leads to detection without redaction (e.g., "KILL" detected but not replaced).
- **Order of replacement:** Using a set makes order unpredictable; no guarantee which term is replaced first.
- **Unicode / accents / normalization:** Some inputs may contain Unicode variants (e.g., zero-widths, diacritics) that break matching.
- **Context matters:** The presence of a word like "suicide" often requires safe-handling (trigger helpers, not just redaction).